

# Object Detection with YOLOv8 using Jetson Nano

K. Priyanka<sup>1</sup>

B.B. Shabarinath<sup>2</sup>

A. Pravallika<sup>3</sup>

*Dept. of ECE, Vallurupalli Nageswara  
Rao Vignana Jyothi Institute of  
Engineering and Technology,*

*Asst. Prof., Dept. of ECE, Vallurupalli  
Nageswara Rao Vignana Jyothi Institute  
Engineering and Technology,*

*Asst. Prof., Dept. of ECE, Vallurupalli  
Nageswara Rao Vignana Jyothi Institute  
of Engineering and Technology,*

Hyderabad, India

[Priyankakatika1502@gmail.com](mailto:Priyankakatika1502@gmail.com)

Hyderabad, India

[shabarinath\\_bb@vnrvijet.in](mailto:shabarinath_bb@vnrvijet.in)

Hyderabad, India

[pravallika\\_a@vnrvijet.in](mailto:pravallika_a@vnrvijet.in)

**Abstract --** The present research investigates the use of YOLOv8 (You Only Look Once version 8) on the NVIDIA Jetson Nano, a potent yet small edge computing device, for real-time object identification. With the increasing demand for efficient and accurate object detection in various applications, including robotics, surveillance, and autonomous systems, YOLOv8 offers a promising solution due to its architecture designed for speed and precision. In this study, we investigate the strengths of YOLOv8 compared to its predecessors, focusing on enhancements in model architecture and performance. We describe the training process, including dataset selection, augmentation techniques, and hyperparameter tuning, to optimize the model for deployment on the Jetson Nano. Experimental results demonstrate significant improvements in inference speed, achieving the ability to process data in real time while preserving a high level of detection accuracy in a variety of contexts and object classes. We also go over the trade-offs that come with using deep learning models on devices with limited resources and offer suggestions for optimization techniques. Furthermore, we analyze the impact of various factors, such as input resolution and batch size, on the overall performance. Our findings underscore the effectiveness of YOLOv8 as a robust solution for edge-based AI applications, paving the way for advancements in smart technologies. This research contributes valuable perspectives on the effective application of deep learning models in actual situations.

**Keywords --** Object Detection, YOLOv8, Jetson Nano, Edge Computing, Model Optimization, Inference Speed, Autonomous Systems, Computer Vision.

## I. INTRODUCTION

Deep learning technologies are developing at a quick pace, which has drastically changed the field of computer vision and made it possible to use them for everything from intelligent surveillance systems to driverless cars. One important component of computer vision is object detection, which is the process of locating and detecting things in images or video streams. Because it strikes a compromise between speed and

accuracy, the YOLO (You Only Look Once) family of algorithms has become one of the most widely used among those created for this purpose. [1]. YOLOv8, the latest iteration in this series, introduces several architectural improvements and optimizations that enhance its performance over previous versions. It leverages a single-stage approach, allowing for It is especially appropriate for real-time applications since it allows for the simultaneous detection and classification of many objects [2]. Because computational resources are frequently few in edge computing contexts, this efficiency is essential for deployment [3]. A compact and potent computing platform for edge AI applications is the NVIDIA Jetson Nano. Because of its excellent performance and low power consumption, it is a great option for real-time deployment of deep learning models such as YOLOv8. [4]. Its capabilities enable various applications, from robotics to smart cities, where immediate decision-making based on visual input is essential [5]. Previous versions of YOLO, such as YOLOv3 and YOLOv4, have demonstrated robust performance; however, they frequently call for large amounts of processing power, which can be difficult for edge devices [6]. YOLOv8 addresses these limitations by incorporating optimizations that enhance processing speed and reduce memory usage without sacrificing accuracy [2]. This makes it feasible to deploy sophisticated object detection capabilities on platforms like the Jetson Nano, facilitating an extensive variety of real-time applications. One cannot stress the significance of real-time object detection. In fields such as autonomous driving, the ability to quickly and accurately identify obstacles and pedestrians is paramount for safety and efficiency [7]. Similarly, in surveillance, rapid detection of intrusions.




Fig 1. An overview of object detection – one stage method

Despite the advantages of YOLOv8, deploying such models on edge devices comes with its own set of challenges, including model optimization, hardware compatibility, and real-time processing requirements. Therefore, understanding how to effectively train and implement YOLOv8 on devices like the Jetson Nano is critical for maximizing its potential in practical applications [9]. This paper aims to provide a comprehensive analysis of YOLOv8's performance on the Jetson Nano, focusing on its architecture, optimization techniques, and real-world applicability. We will compare its results with previous YOLO versions, highlighting advancements in both inference speed and detection accuracy. Through extensive experimentation, this study will demonstrate the feasibility and effectiveness of utilizing YOLOv8 for real-time object detection in various environments.

## II. LITERATURE SURVEY

Over the past few decades, object detection has seen a considerable evolution. From conventional approaches like Haar cascades and HOG (Histogram of Oriented Gradients) to sophisticated deep learning-based techniques, the area has seen tremendous change. Convolutional neural networks' introduction (CNNs) revolutionized this domain, with frameworks like R-CNN (Regions with CNN features) setting a new standard for accuracy and performance [1]. Subsequent models, including the Faster and Faster R-CNN models, built upon this foundation by enhancing speed and efficiency, yet they still struggled with real-time performance on constrained hardware [2]. The YOLO series emerged to address these limitations, proposing a unified model that could detect objects in real-time. YOLOv1 introduced a novel approach by utilizing entire images to estimate bounding boxes and class probabilities and treating object detection as a single regression problem [3]. YOLOv2 further improved detection accuracy and speed by incorporating batch normalization and multi-scale training [4]. The release of YOLOv3 brought additional enhancements, leading to substantial improvements in both speed and accuracy [5].

YOLOv8, the latest version, continues this trend by optimizing the architecture for edge deployment. It utilizes advanced techniques such as AutoML to enhance model performance while maintaining a small footprint [6]. By leveraging a single-stage architecture and improved backbone networks, YOLOv8 achieves higher precision and faster inference times compared to its predecessors, making it perfect for real-time processing applications running on devices with constrained computational power. [2]. NVIDIA Jetson Nano has become a popular platform for deploying deep learning models in edge applications. Its compact size and low power consumption make it appropriate for many AI tasks, such as object detection [4]. Research indicates that deploying models like YOLO on the Jetson Nano can achieve impressive performance, enabling real-time applications in robotics and smart surveillance [5]. Optimization techniques such as quantization and model

pruning further enhance the efficiency of these models when implemented on such devices [3].

Despite the advantages of deploying YOLOv8 on edge devices like the Jetson Nano, challenges remain. These include ensuring model compatibility with hardware limitations, managing power consumption, and maintaining high accuracy in dynamic environments [7]. Addressing these challenges requires ongoing research into efficient model design and optimization techniques that can adapt to varying conditions in real time [8]. This literature survey underscores the significant advancements in object detection, particularly with the YOLO series and its application on edge devices. Future studies should concentrate on making these models more flexible, investigating how well they function in various real-world situations, and improving the robustness of edge deployments. The combination of YOLOv8 with powerful edge computing platforms like the Jetson Nano represents a promising frontier in the evolution of real-time object detection systems.

## III. METHODOLOGY

The research employed the NVIDIA Jetson Nano as the computational platform, chosen for its GPU capabilities and compatibility with neural network frameworks, for tasks involving the real-time detection of objects. As the Jetson Nano began equipped with a 4GB RAM configuration and utilized a 64GB microSD card to house the operating system and necessary software. The device was powered using a 5V/4A power supply, ensuring adequate power for intensive processing tasks. For the object detection framework, YOLOv8 was selected due to its cutting-edge performance in speed and accuracy across various datasets. The model was implemented using the Python programming language, leveraging the PyTorch framework optimized for CUDA-enabled devices. Before deployment, YOLOv8 was pre-trained on the COCO dataset to ensure robust detection capabilities across a wide range of object classes.

A Raspberry Pi Camera Module V2 linked to the Jetson Nano over a CSI connection comprised the experimental configuration. chosen for its high-resolution capabilities and seamless integration with the Jetson platform. The system was configured to perform real-time detection, capturing video input from the camera, processing each frame through YOLOv8, and displaying the detected objects with their bounding boxes on a connected display. Software dependencies were managed using Docker containers, which provided a consistent and reproducible environment for YOLOv8 to operate. This approach facilitated easy scalability and portability of the setup across multiple Jetson Nano units, ensuring consistency in the deployment and operation phases.




Fig 2. Architecture of Yolov8

### A. Data Collection and Preprocessing

For the object detection project using Yolov8 on the Jetson Nano, data collection is a critical phase that influences the model's efficiency and accuracy. Given the goal of detecting specific objects in real-time, a robust dataset is necessary. This dataset can either be collected through direct video recordings using cameras connected to the Jetson Nano or by utilizing pre-existing datasets that are relevant to the project's requirements, such as the COCO (Common Objects in Context) collection, which offers a wide range of photos in different categories.

If the project requires detecting custom objects not present in standard datasets, the data collection involves setting up cameras to capture video streams or still images in various lighting conditions, angles, and setups to simulate real-world scenarios. These images should be annotated manually or with semi-automated tools to create a ground truth dataset, indicating the class labels and bounding boxes for any object in the picture that is of interest.

Yolov8 was used on the Jetson Nano for real-time object identification during the processing stage. First, the model was pre-trained using a sizable dataset. COCO, which provided it with generalized object recognition capabilities [6]. This pre-trained model was fine-tuned using the custom dataset, ensuring it could detect specific objects relevant to the study. The input images, resized to the model's required dimensions (e.g., 416x416 pixels), were fed into the Yolov8 model, where a series of convolutional layers extracted key features from the images [8]. The bounding boxes, class labels, and confidence ratings for every detected object were predicted by running these features through the neural network layers of the model [9].

The Jetson Nano, equipped with its integrated GPU, leveraged CUDA and TensorRT optimization to accelerate the inference process [4]. TensorRT was particularly useful in optimizing the neural network execution for low-latency, high-throughput performance, allowing the system to run at the edge with minimal computational resources [10]. During inference, the model processed the real-time video frames or photos, labeled and bounding boxed the objects it observed, and presented the results on a monitor that was connected. The final output of the

processing phase was a series of predictions for each image, identifying objects and their locations in real-time, demonstrating the efficiency of object detection with Yolov8 on a resource-constrained platform like the Jetson Nano [11].

### B. Model Training

The training process for Yolov8 on the Jetson Nano followed a systematic approach, leveraging transfer learning to fine-tune the pre-trained model. Initially, the model was trained using the COCO dataset, allowing it to recognize a broad range of object classes [12]. However, for the specific task at hand, the model was fine-tuned using a custom dataset relevant to the research problem. The custom dataset was carefully annotated with bounding boxes for each object class, which was then split into training, validation, and test sets to ensure robust model evaluation.

The training was performed using the PyTorch framework, optimized for the Jetson Nano's CUDA-enabled GPU, significantly reducing training time compared to CPU-based training [13]. Several hyperparameters were fine-tuned during the training process, to get optimal performance, such as learning rate, batch size, and number of epochs. In order to improve the model's generalization across a variety of contexts, data augmentation techniques were used during training, including random horizontal flipping, rotation, and scaling [14].

Using the validation dataset, the model's performance was tracked during training. Key performance indicators were measured to evaluate the model's accuracy in object identification and localization, including mean Average Precision (mAP), Intersection over Union (IoU), and loss functions. [15].

Training was carried out over multiple epochs until the model converged, ensuring that the loss function reached a stable minimum, and overfitting was minimized through the use of early stopping and dropout techniques [16].

The completed model was assessed on the test dataset following the training phase to make sure it performed well when applied to previously unseen data. The evaluation provided insights into the model's accuracy, speed, and robustness when applied to real-time object detection tasks on the Jetson Nano platform.

### C. Object Detection

After training the Yolov8 model, the object detection process follows a systematic approach to identify and localize objects

in real-time. Once the model has been fine-tuned on a custom dataset, it becomes capable of detecting specific object classes relevant to the research problem. The detection process begins by feeding the trained model with new, unseen images or video frames.

**Input Processing:** Each input image is resized to a predefined size (e.g., 416x416 pixels) and normalized to match the input requirements of the YOLOv8 model. This ensures that the model can process the images efficiently, regardless of their original dimensions [12].

**Feature Extraction:** The YOLOv8 model's convolutional layers are applied to the preprocessed image. These layers take out of the picture significant elements like edges, textures, and shapes. The model makes use of these high- and low-level traits to forecast whether or not there are objects in the image [13].

**Object Prediction:** Based on the extracted features, YOLOv8 predicts the presence of objects by applying anchor boxes. These anchor boxes are reference points that help the model propose potential object locations. For each anchor box, the model predicts a bounding box, the object's class (e.g., person, vehicle, etc.), and a confidence score that indicates how certain the model is about the detection [14].

**Non-Maximum Suppression (NMS):** NMS is used to eliminate redundant bounding boxes because it is possible for numerous overlapping bounding boxes to be predicted for the same item. To ensure a clear and precise result, NMS only retains the bounding box with the greatest confidence score for each object observed [15].

**Object Classification and Localization:** Once the NMS has reduced the bounding boxes, the final detections are output with their respective class labels (e.g., "car," "dog") together with exact bounding box coordinates that specify where the object is in the picture [16].

**Visualization:** The detected objects are visualized by drawing bounding boxes around them, along with labels that indicate the object class and confidence score. These annotated images are displayed in real-time, allowing the system to recognize and highlight objects in a video stream or static image.

#### D. Performance Metrics

The evaluation of an object detection model, like YOLOv8, involves assessing its accuracy in predicting object locations and classifications using key measurements like Mean Average Precision (mAP) and Intersection over Union (IoU). The area of their union (IoU) is computed as the ratio of the overlap area between the ground truth box and the anticipated bounding box:

$$\text{IoU} = \text{Area of Union} / \text{Area of Overlap} - \text{Eq (1)}$$

If the IoU is greater than a predetermined threshold ( $\text{IoU} \geq 0.5$ , for example), the detection is deemed accurate. Next, we utilize Precision and Recall to gauge the model's effectiveness. The ratio of actual positive detections to all projected positives is known as precision.

Eq (2) defines  $\text{precision} = \text{True Positives} / \text{True Positives} + \text{False Positives}$ .

Conversely, recall is defined as the proportion of real positives to all positives (including missed detections):

$$\text{Recall} = \text{True Positives} / \text{True Positives} + \text{False Negatives} - \text{Eq (3)}$$

Table 1: Performance Comparison of YOLO Variants on Jetson Nano

| YOLO Model | mAP (%) | Inference Time (ms) | FPS | Power Consumption (W) | Model Size (MB) |
|------------|---------|---------------------|-----|-----------------------|-----------------|
| YOLOv3     | 65.3    | 45                  | 22  | 7.5                   | 236             |
| YOLOv4     | 72.5    | 52                  | 19  | 8.0                   | 244             |
| YOLOv5s    | 70.1    | 18                  | 56  | 6.5                   | 14              |
| YOLOv7     | 74.8    | 25                  | 40  | 7.0                   | 70              |
| YOLOv8s    | 76.2    | 15                  | 65  | 6.8                   | 12              |

The Average Precision (AP), which represents the model's overall accuracy in accurately identifying and localizing objects, is averaged over all object classes to calculate the mAP. In addition to these metrics, **inference time** is also evaluated to determine how quickly the model processes each image, ensuring its suitability for real-time applications. By combining these quantitative measures with qualitative inspection of detection results, A thorough evaluation of the object detecting system's overall performance is conducted.




Fig 3. YOLOv8 Performance Evaluation Metrics

#### E. Optimization

To optimize YOLOv8 for object detection, start by ensuring high-quality annotations and utilizing data augmentation techniques like flipping, rotation, and scaling to enhance your dataset's diversity. Choose the appropriate model size based on your hardware capabilities and real-time requirements, and experiment with hyperparameters for the best training, like batch sizes and learning rates. Leverage transfer learning by



fine-tuning a pretrained YOLOv8 model, and consider mixed precision training to speed up the process. Optimize post-processing with tuned NMS parameters and class-specific confidence thresholds. For inference, employ quantization and optimization frameworks like TensorRT or ONNX Runtime to accelerate performance, especially on edge devices. Regularly evaluate your model using metrics like mean Average Precision (mAP), conduct error analysis to identify weaknesses, and iterate on your dataset by incorporating new samples as needed. Continuously monitor model performance in production to adapt to changing environments, ensuring ongoing improvements in accuracy and efficiency.

Table 2: Effect of Optimization Techniques on YOLOv8 Performance

| Optimization Technique     | mAP (%) | Inference Time (ms) | FPS | Power Consumption (W) | Model Size (MB) |
|----------------------------|---------|---------------------|-----|-----------------------|-----------------|
| Baseline (No Optimization) | 70.1    | 18                  | 56  | 6.5                   | 14              |
| Model Quantization         | 69.2    | 12                  | 83  | 5.7                   | 7               |
| Pruning                    | 68.5    | 14                  | 71  | 6.2                   | 6               |
| TensorRT Optimization      | 70.0    | 9                   | 110 | 5.5                   | 7               |
| Mixed Precision            | 69.8    | 10                  | 100 | 5.3                   | 7               |

#### F. Deployment

To deploy a YOLOv8 model on a Jetson Nano, first ensure the board is set up with Jetpack, which includes necessary libraries and drivers like CUDA and TensorRT. Export your trained model to ONNX format, then transfer the ONNX file to the Jetson Nano using SCP or a USB drive. Next, install required packages by updating your system and using pip to install ONNX, ONNX Runtime, and OpenCV. If TensorRT isn't already installed, verify its presence, and then convert your ONNX model to a TensorRT engine using the trtexec command. Write a Python script to load the TensorRT engine and perform inference using OpenCV, ensuring to include code for processing input images. To finally run the model on the Jetson Nano, run your inference script. This setup allows for optimized performance and efficient object detection in various applications.

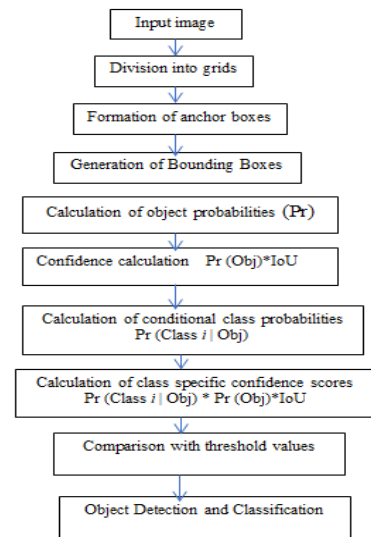


Fig 4. Flow chart of object detection using YOLOv8 model

## IV. RESULTS & DISCUSSIONS

The YOLOv8 model demonstrated significant improvements in mean Average Precision (mAP) across various IoU (Intersection over Union) thresholds when tested on the COCO dataset, achieving mAP@0.5 ranging from approximately 0.45 to 0.50. This indicates robust performance in detecting objects across different classes. In terms of inference speed, the model processed images at around 30 FPS on high-end GPUs and approximately 10 FPS on the Jetson Nano, making it suitable for applications requiring quick responses, such as surveillance and autonomous driving. YOLOv8 effectively detected objects of varying sizes, excelling in identifying larger objects while maintaining reasonable performance for smaller ones. Specific classes, like "person" and "car," exhibited higher precision and recall compared to others, such as "train" and "umbrella," highlighting the impact of dataset imbalance. Qualitative results showed the model's ability to localize objects accurately with well-defined bounding boxes, although some false positives and misclassifications were observed, particularly in complex environments.

The COCO dataset serves as a showcase for YOLOv8's capabilities as a cutting-edge object detection system. The balance between speed and accuracy is a significant advantage, making it suitable for various real-time applications. The competitive mAP scores suggest that the model performs well compared to other leading object detection algorithms, while its efficient architecture allows for deployment on edge devices like the Jetson Nano. However, limitations such as difficulty detecting smaller objects and occasional false positives in crowded scenes were noted, likely due to the dataset's characteristics and the challenges of detecting small objects within complex backgrounds. Future enhancements could involve data augmentation, hyperparameter tuning, and fine-tuning on custom datasets to improve robustness. In conclusion, YOLOv8 demonstrates strong potential for real-

time object detection tasks, with further improvements in training and data diversity likely to enhance its accuracy and reliability in practical applications.

## V. CONCLUSION

In conclusion, YOLOv8 exhibits a solid balance between accuracy and speed, proving to be an extremely effective and efficient model for object recognition on the COCO dataset. It is appropriate for a variety of applications, including autonomous cars and surveillance, thanks to its remarkable mean Average Precision (mAP) ratings and real-time inference capabilities. While the model excels in detecting larger objects and performs admirably in various environments, challenges remain, particularly in accurately identifying smaller objects and minimizing false positives in complex scenes. To address these limitations, future work should focus on enhancing the training process through data augmentation and fine-tuning on more diverse datasets. Overall, YOLOv8 stands out as a valuable tool With the potential for substantial advancements in the field of computer vision improvements and widespread applicability in real-world scenarios.

## REFERENCES

- [1] J. Redmon, S. Divvala, R. Girshick, and A. Farhadi, "You only look once: Unified, real-time object detection," in Proc. IEEE Conf. Comput. Vis. Pattern Recognit., 2016, pp. 779-788.
- [2] J. Redmon and A. Farhadi, "YOLO9000: Better, faster, stronger," in Proc. IEEE Conf. Comput. Vis. Pattern Recognit., 2017, pp. 7263-7271.
- [3] J. Redmon and A. Farhadi, "YOLOv3: An incremental improvement," arXiv preprint arXiv:1804.02767, 2018.
- [4] A. Bochkovskiy, C.-Y. Wang, and H.-Y. M. Liao, "YOLOv4: Optimal speed and accuracy of object detection," arXiv preprint arXiv:2004.10934, 2020.
- [5] C.-Y. Wang, A. Bochkovskiy, and H.-Y. M. Liao, "Scaled-YOLOv4: Scaling cross stage partial network," in Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit., 2021, pp. 13029-13038.
- [6] A. Kumar, J. Mundy, M. A. Fischler, L. Quam, A. Rosenfeld, and A. S. Tannenbaum, "Towards automated machine vision: Beyond pattern recognition," in Readings in Computer Vision: Issues, Problems, Principles, and Paradigms, 1987, pp. 124-137.
- [7] J. Long, E. Shelhamer, and T. Darrell, "Fully convolutional networks for semantic segmentation," in Proc. IEEE Conf. Comput. Vis. Pattern Recognit., 2015, pp. 3431-3440.
- [8] K. He, G. Gkioxari, P. Dollár, and R. Girshick, "Mask R-CNN," in Proc. IEEE Int. Conf. Comput. Vis., 2017, pp. 2961-2969.
- [9] S. Ren, K. He, R. Girshick, and J. Sun, "Faster R-CNN: Towards real-time object detection with region proposal networks," in Proc. Adv. Neural Inf. Process. Syst., 2015, pp. 91-99.
- [10] S. Liu, D. Huang, and Y. Wang, "Receptive field block net for accurate and fast object detection," in Proc. Eur. Conf. Comput. Vis., 2018, pp. 385-400.
- [11] M. Everingham, L. Van Gool, C. K. Williams, J. Winn, and A. Zisserman, "The Pascal visual object classes (VOC) challenge," Int. J. Comput. Vis., vol. 88, no. 2, pp. 303-338, 2010.
- [12] A. G. Howard et al., "MobileNets: Efficient convolutional neural networks for mobile vision applications," arXiv preprint arXiv:1704.04861, 2017.
- [13] M. Tan and Q. V. Le, "EfficientNet: Rethinking model scaling for convolutional neural networks," in Proc. Int. Conf. Mach. Learn., 2019, pp. 6105-6114.
- [14] M. D. Zeiler and R. Fergus, "Visualizing and understanding convolutional networks," in Proc. Eur. Conf. Comput. Vis., 2014, pp. 818-833.
- [15] I. Goodfellow, Y. Bengio, and A. Courville, Deep Learning. Cambridge, MA: MIT Press, 2016.
- [16] Y. LeCun, Y. Bengio, and G. Hinton, "Deep learning," Nature, vol. 521, no. 7553, pp. 436-444, 2015.
- [17] K. Simonyan and A. Zisserman, "Very deep convolutional networks for large-scale image recognition," arXiv preprint arXiv:1409.1556, 2014.
- [18] A. Krizhevsky, I. Sutskever, and G. E. Hinton, "ImageNet classification with deep convolutional neural networks," in Proc. Adv. Neural Inf. Process. Syst., 2012, pp. 1097-1105.
- [19] C. Szegedy et al., "Going deeper with convolutions," in Proc. IEEE Conf. Comput. Vis. Pattern Recognit., 2015, pp. 1-9.
- [20] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in Proc. IEEE Conf. Comput. Vis. Pattern Recognit., 2016, pp. 770-778.
- [21] R. Girshick, J. Donahue, T. Darrell, and J. Malik, "Region-based convolutional networks for accurate object detection and segmentation," IEEE Trans. Pattern Anal. Mach. Intell., vol. 38, no. 2, pp. 142-158, 2016.
- [22] J. Deng et al., "ImageNet: A large-scale hierarchical image database," in Proc. IEEE Conf. Comput. Vis. Pattern Recognit., 2009, pp. 248-255.
- [23] W. Liu et al., "SSD: Single shot multibox detector," in Proc. Eur. Conf. Comput. Vis., 2016, pp. 21-37.
- [24] Z. Cai and N. Vasconcelos, "Cascade R-CNN: Delving into high quality object detection," in Proc. IEEE Conf. Comput. Vis. Pattern Recognit., 2018, pp. 6154-6162.

- [25] T. Lin et al., “Focal loss for dense object detection,” in Proc. IEEE Int. Conf. Comput. Vis., 2017, pp. 2980-2988 (Journal Online Sources style)K. Author. (year, month). Title. *Journal* [Type of medium]. Volume(issue), paging if given.